



Tetris: Predictive Pod Placement Strategy for Kubernetes

Vasudev Bongale (vabongal)

Sohail Shaikh (sashaikh)

Kyle Martin (kdmarti2)

26-April-2019



Agenda

- Problem Statement
- Idea
- Architecture
- Forecasting Strategy
 - Results
- Tetris-scheduler
 - Results
- Threat To Validity
- Limitations
- Future Work

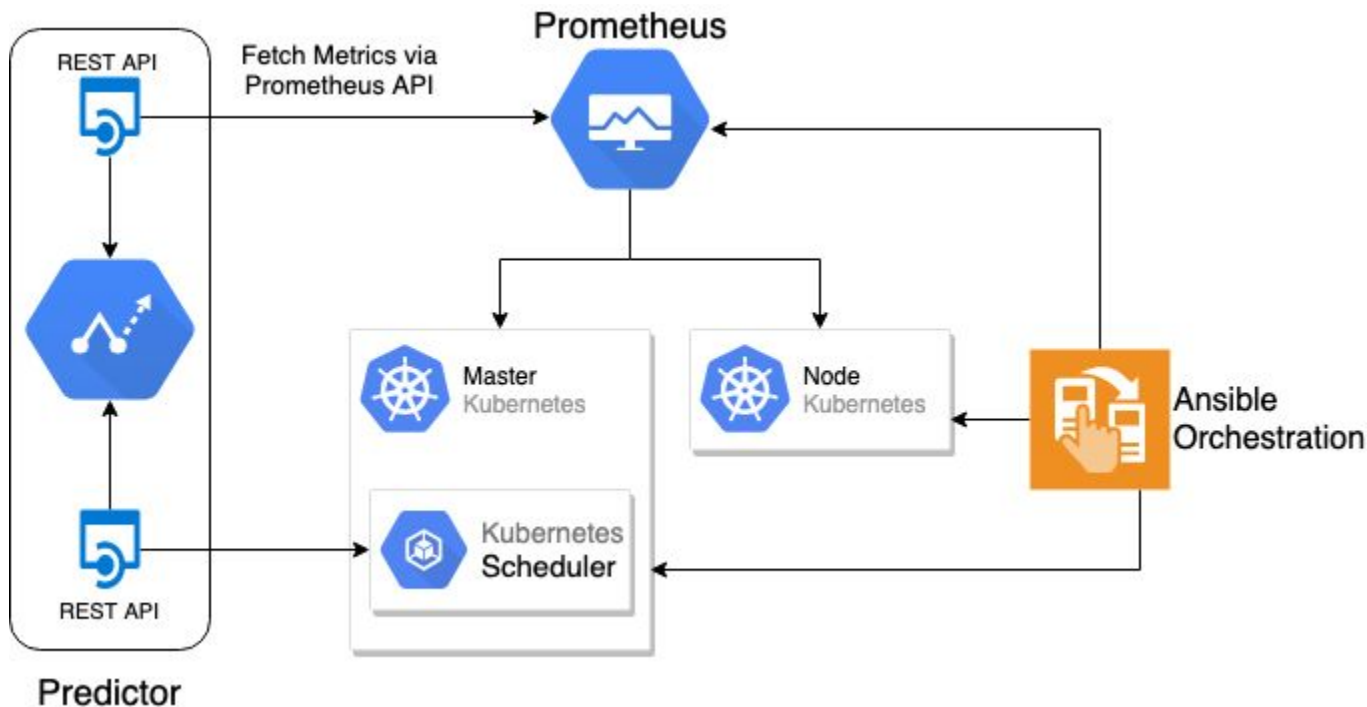
Problem Statement

Kubernetes does not consider I/O utilization when provisioning pods. This could lead to performance degradation on an I/O bottleneck K8s Node. We propose to introduce a predictive pod placement strategy to avoid resource bottlenecks.

Idea

- Schedule based on resource requirements.
 - Manually labelling pod by workload type.
 - Memory/IO intensive.
- Ensemble of prediction algorithms.
 - Intelligent placement.
 - Choose a node for a pod.

Architecture



Training Data Collection

1. Three different K8s stress-ng deployments.

cpu	<code>stress-ng --cpu \$workers --cpu-load \$cpu_load --cpu-load-slice \$load_slice -t \$time</code>
mem	<code>/usr/bin/stress-ng --vm \$num_workers --vm-bytes 40% --vm-method rand-set --t \$time_out</code>
io	<code>stress-ng --iomix \$iomix --iomix-bytes 1g -t \$timeout</code>

2. Stats exported by Prometheus Node Exporter.

Prometheus Queries

- CPU Usage (%): `"100 - (avg by (instance) (irate(node_cpu{job='k8s-nodes',mode='idle'}[2m])) * 100)"`
- Used Memory (%): `"100 * (1 - ((node_memory_MemFree{job='k8s-nodes'} + node_memory_Cached{job='k8s-nodes'} + node_memory_Buffers{job='k8s-nodes'}) / node_memory_MemTotal{job='k8s-nodes'}))"`
- Disk I/O utilization (%): `"100 * (rate(node_disk_io_time_ms{device='xvda', job='k8s-nodes'}[5m])/1000 or irate(node_disk_io_time_ms{device='xvda', job='k8s-nodes'}[5m])/1000)"`

APIs

- Components communicate using REST APIs.
- Prometheus Module
 - Get metrics from prometheus-server
 - Prometheus API
 - GET /api/v1/query
 - GET /api/v1/query_range
- Scheduler Queries a Flask API to interact with the module.

Prediction Module

- Online Training
- Ensemble of prediction algorithms:
 - ARIMA
 - Holt Winters
 - WMA
- Training period is last 6 hours.

Precompute Module

- Runs every 5 minutes to check the accuracies of models for cpu, io and memory utilizations.
- Picks the model with the least mean error percent and that model is used for prediction within that window.

Auto Regressive and Integrated Moving Average

- Good for detecting seasonality.
- Expects a stationary series.
- Key concepts are order, lags and differencing:

Order: AR(p)

$$Y_t = \beta_0 + \beta_1 Y_{t-1} + \beta_2 Y_{t-2} + \dots + \beta_p Y_{t-p} + \epsilon_t$$

ARMA(p, q) : q-> lag versions of forecast errors

$$Y_t = \beta_0 + \beta_1 Y_{t-1} + \beta_2 Y_{t-2} + \dots + \beta_p Y_{t-p} \\ + \epsilon_t + \theta_1 \epsilon_{t-1} + \dots + \theta_q \epsilon_{t-q}$$

ARIMA(p, d, q):

d = 0: no differencing

d = 1: performing differencing once

Apply differencing to remove trends from the series.

Tuning ARIMA parameters:

- Choosing parameters (p , d , q) is difficult.
- Visual inspection is required to figure out seasonality and trend.
- We use AIC values as a criterion for choosing best model using a grid searching technique.

Holt Winter's Exponential Smoothing

- Can detect trends as well as seasonality.
- Popular because it's fast and cheap to compute.
- Key concepts: smoothing constants and updating equations.

Forecast = estimated level + trend + seasonality at most recent time point

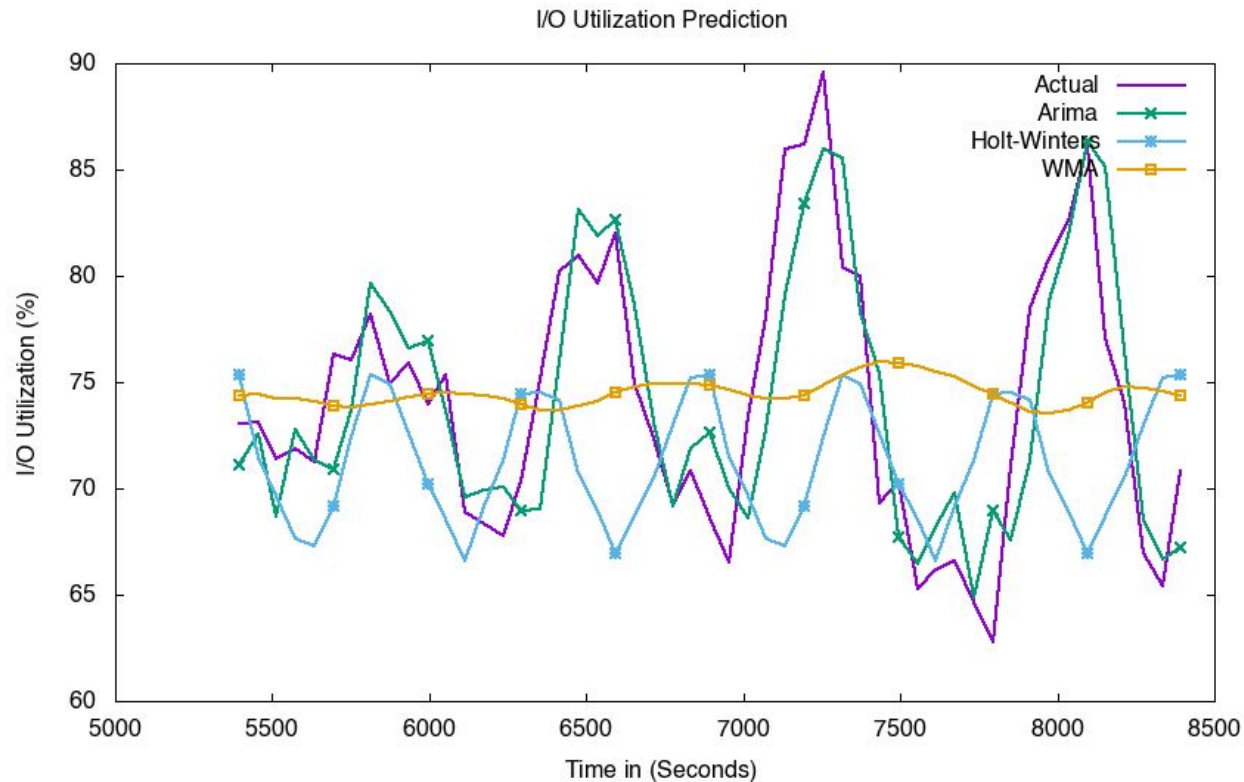
$$F_{t+k} = L_t + kT_t + S_{t+k-M}$$

Level(L_t) *trend*(T_t) *Seasonality*(S_t) with M seasons


Prediction Algorithm Analysis Methodology

1. Workload: I/O, Memory and CPU stressors (stress-ng).
2. Sampling rate: 10s
3. Training period: 6 hrs
4. Forecasting window: 6 points in the future (1 minute).
5. Shift prediction window by a minute.
6. Repeat 4, 5 for 1 hour.
7. This analysis was carried out in the offline mode.

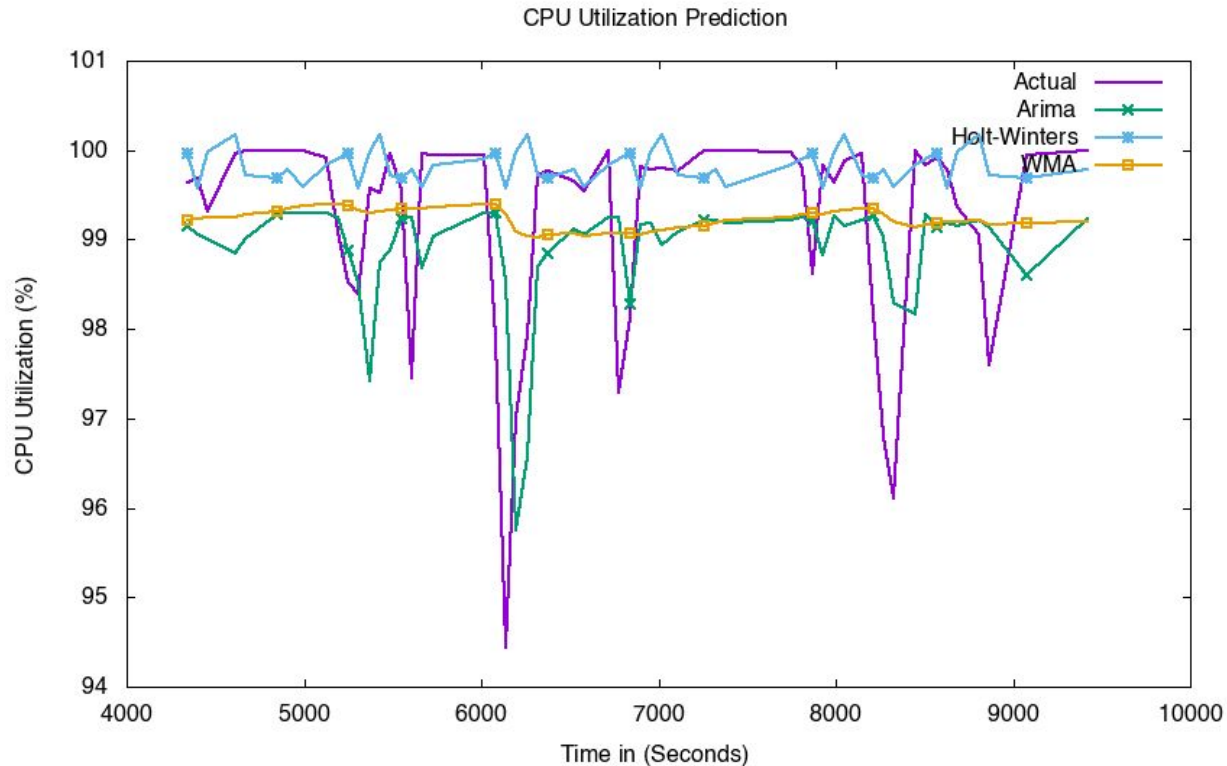
Disk IO% Utilization



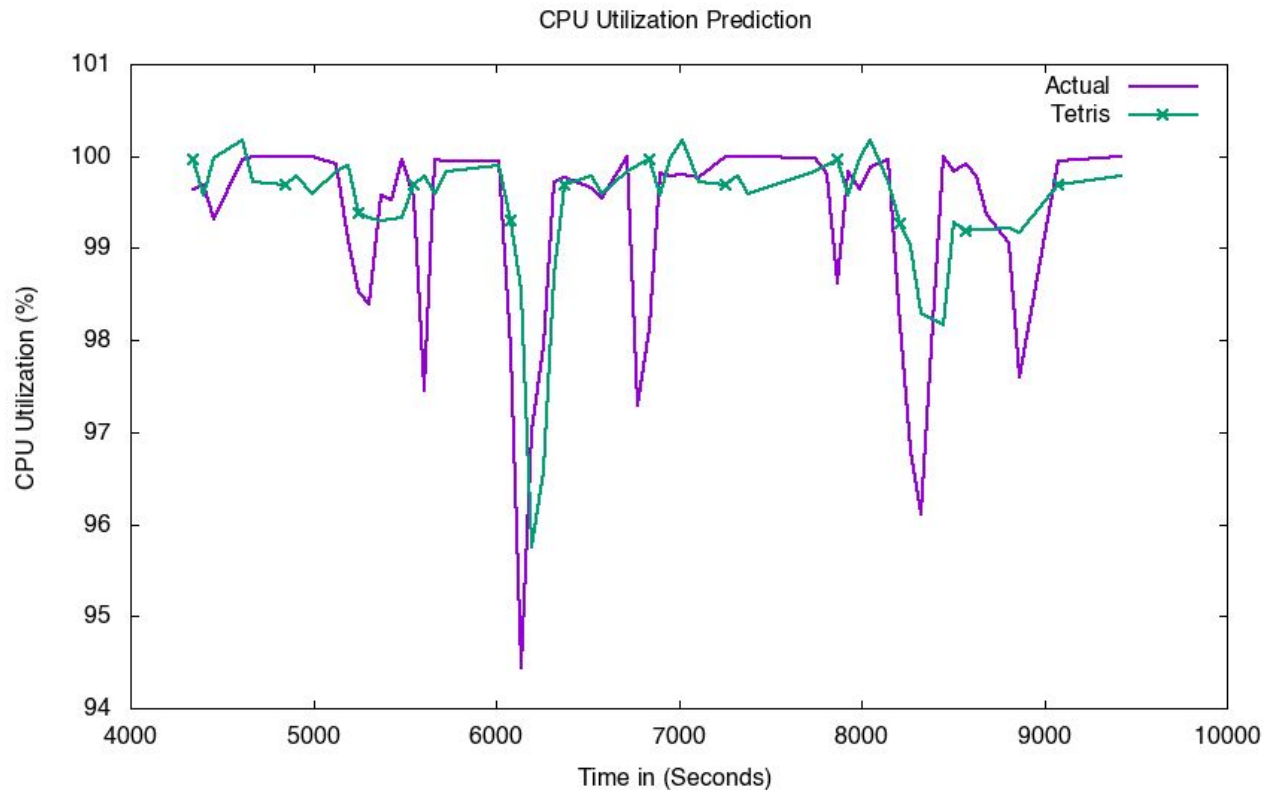
Disk I/O Utilization Prediction using Tetris



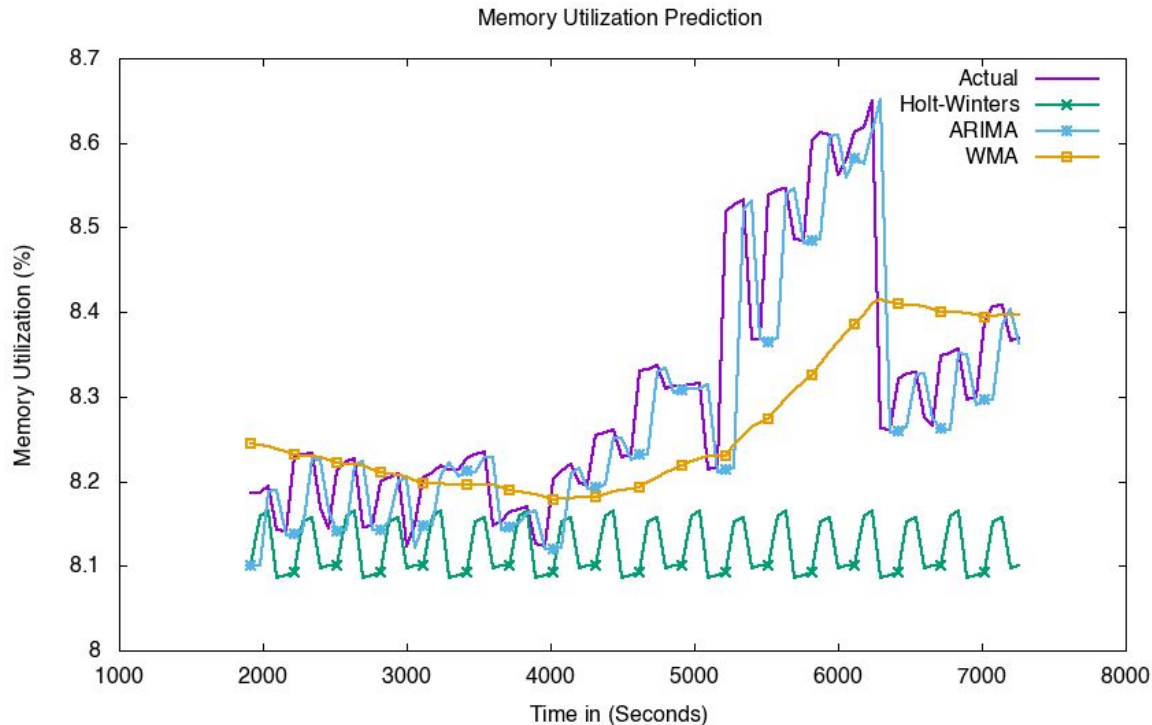
CPU Usage % Prediction

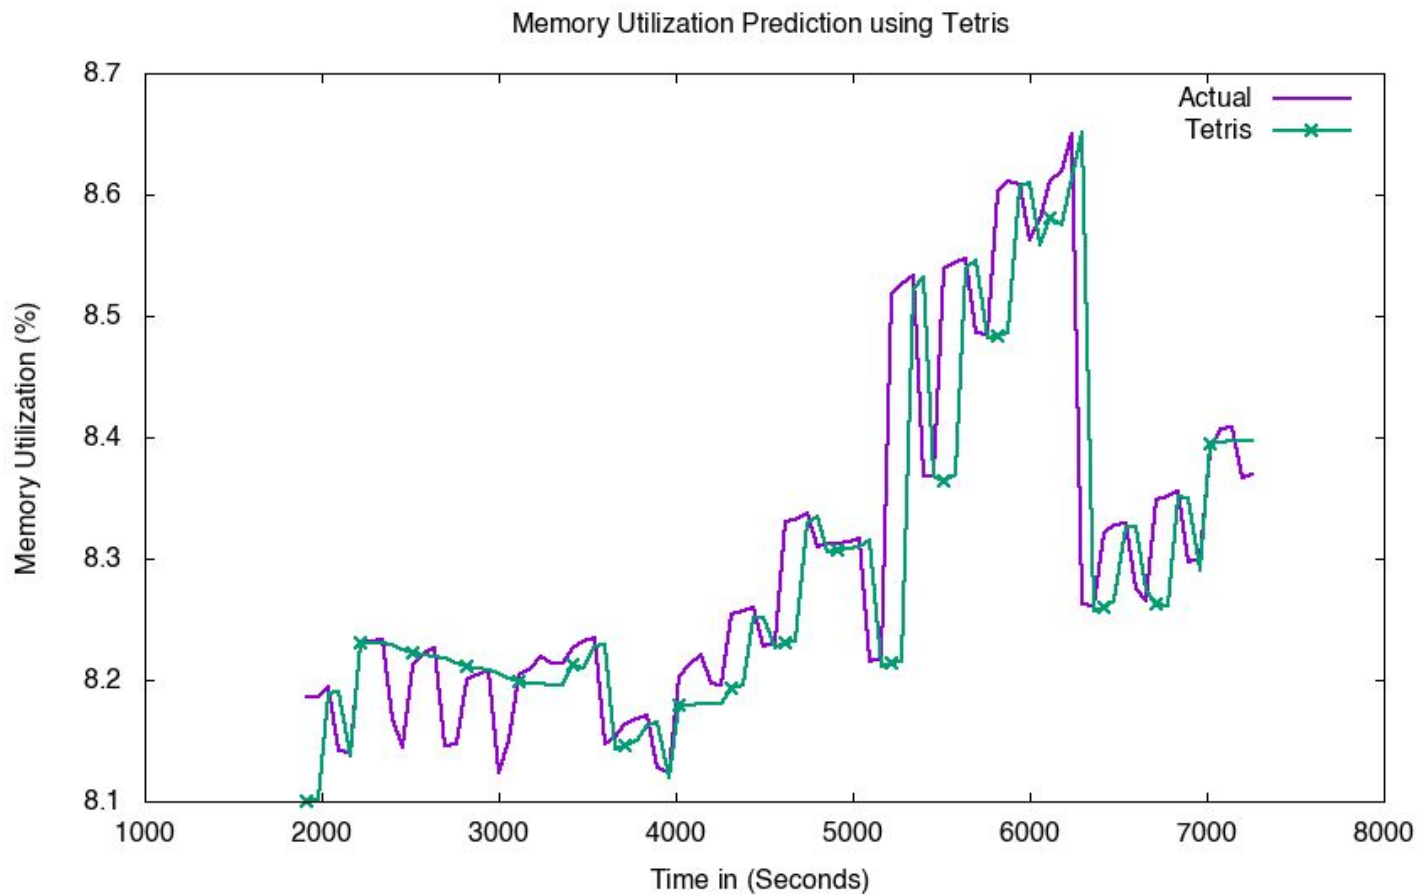


CPU Utilization Prediction using Tetris



Memory Usage %






Prediction Mean Error Percent


Algorithm	Disk I/O Utilization Accuracy	Memory Utilization Accuracy	CPU Utilization Accuracy
ARIMA	3.5%	0.64%	1%
Holt Winters	7.2%	2.1%	0.8%
WMA	6.6%	1.05%	0.9%
*Tetris	3.65%	0.55%	0.7%

Tetris Scheduler



Selecting A Scheduler

- Tetris-scheduler-io
- Tetris-scheduler-mem
- Tetris-scheduler-cpu

```
admin/sched/pod2.yaml   
apiVersion: v1  
kind: Pod  
metadata:  
  name: annotation-default-scheduler  
  labels:  
    name: multischeduler-example  
spec:  
  schedulerName: default-scheduler  
  containers:  
  - name: pod-with-default-annotation-container  
    image: k8s.gcr.io/pause:2.0
```

```
spec:  
  schedulerName: tetris-scheduler-io  
  containers:  
  - image: mysql:5.6  
    name: mysql  
    env:
```


SLO MySql Response Time Methodology

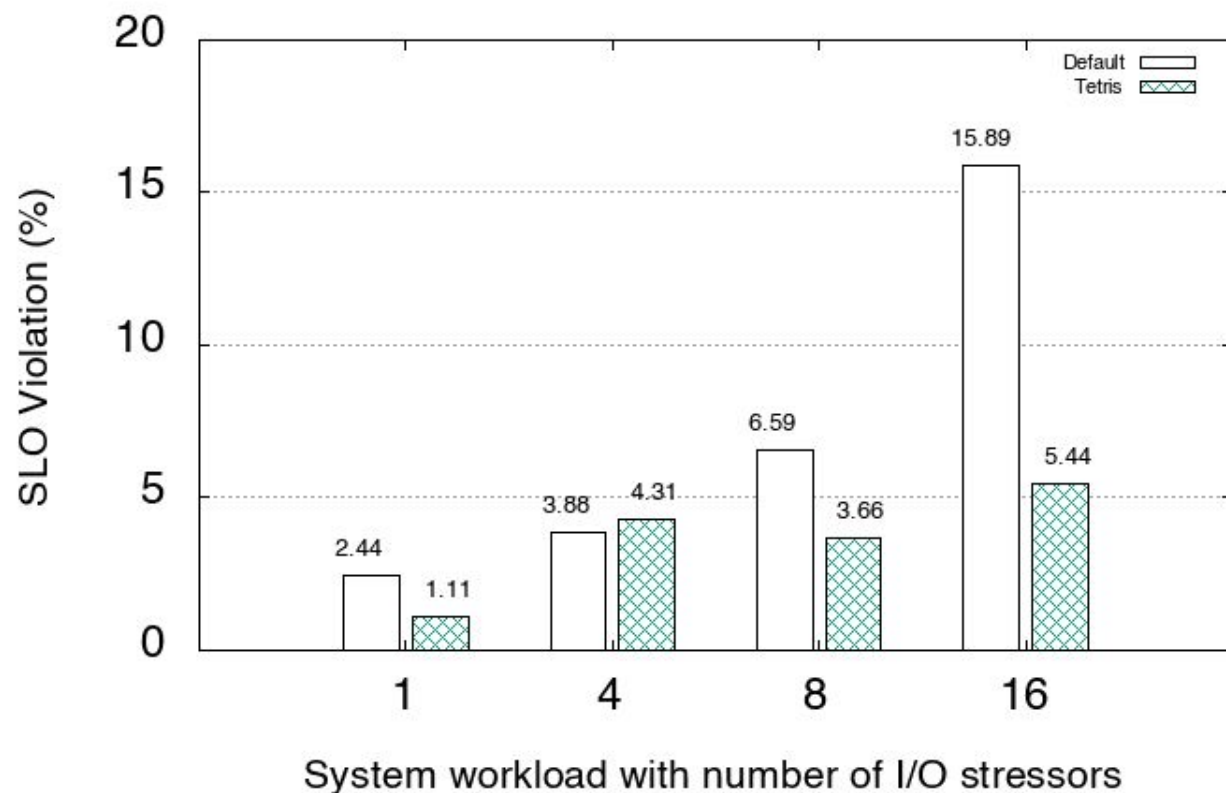
1. Deploy K8s MySQL pod.
2. Record response time of 10K insert requests (write).
3. Record response time of 10K query requests (read).
4. Determine 99th percentile respectively.

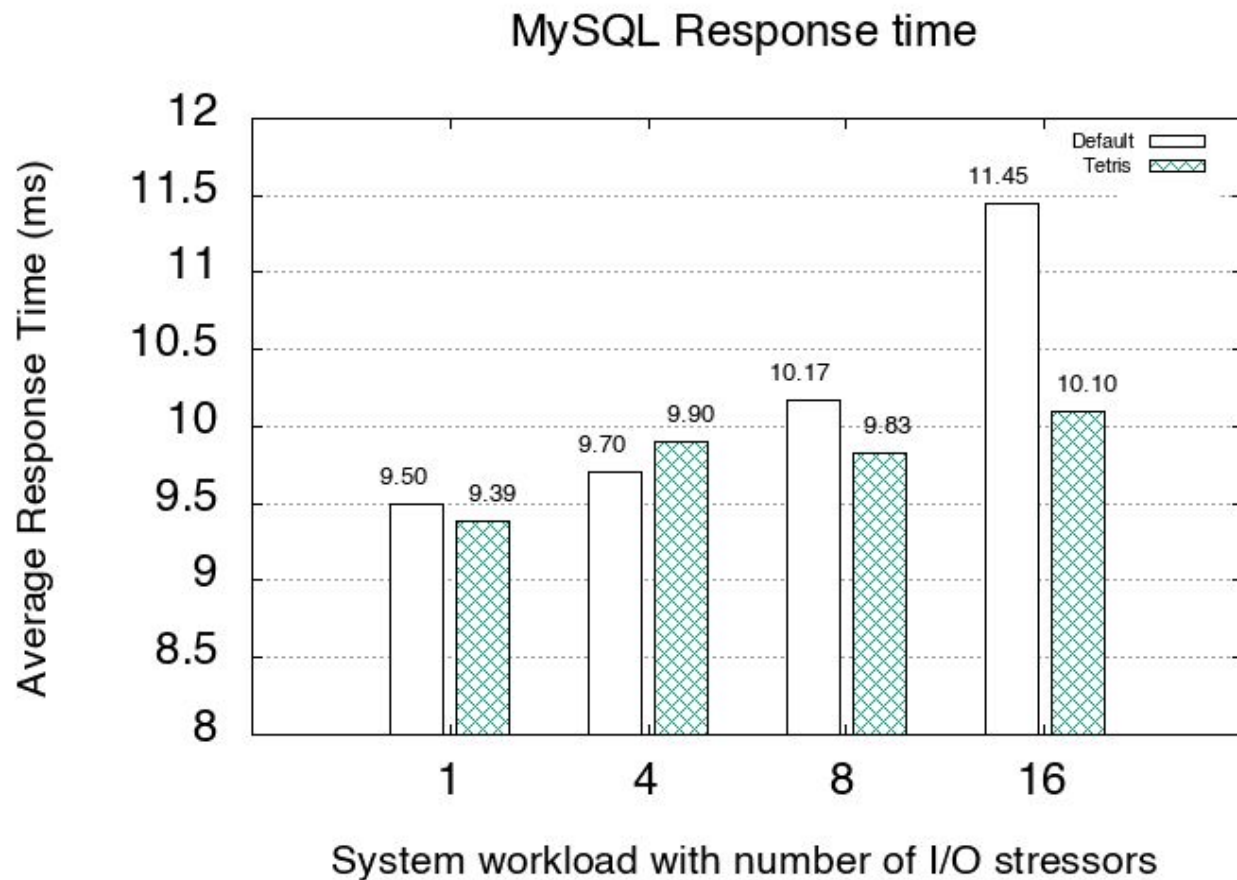
Write	11.84 ms
Read	28.47 ms

IO Aware System Evaluation Methodology

1. Deploy io-stressors.
2. Deploy MySql K8s deployment.
3. Record response time of 10K insert requests.
4. Record response time of 10K query requests.
5. Repeat 1-4 for K8s scheduler and Tetris.

MySQL Write SLO Violations

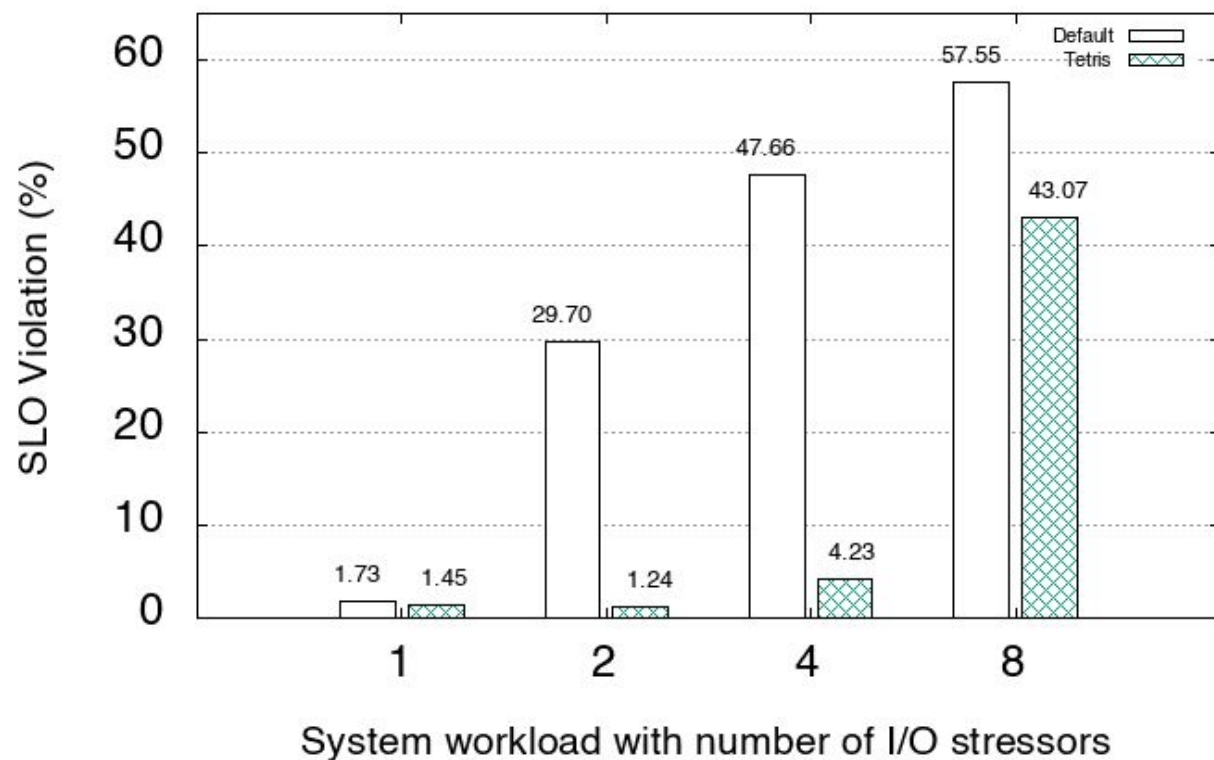





Memory Aware System Evaluation Methodology

1. Deploy mem-stressors.
2. Deploy MySql K8s deployment.
3. Record response time of 10K insert requests.
4. Record response time of 10K query requests.
5. Repeat 1-4 for K8s scheduler and Tetris.

MySQL Write SLO Violations





Threat to Validity

- CPU credits in AWS - allows burstable performance*
- Data Influence of past experiments
 - Data is not representative of actual workload.

* <https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/burstable-performance-instances.html>

Limitations

- Does not react well with fast pod placement requests.
- If the workload does not follow a pattern prediction node placement may not be accurate.

Future work

- Consolidation of containers based on different workloads.
 - Leveraging PAC paper for pod signature extraction.
- Background Grid searching to optimize algorithms.

Questions?